




KSQL - Introduction

Heikki Nousiainen

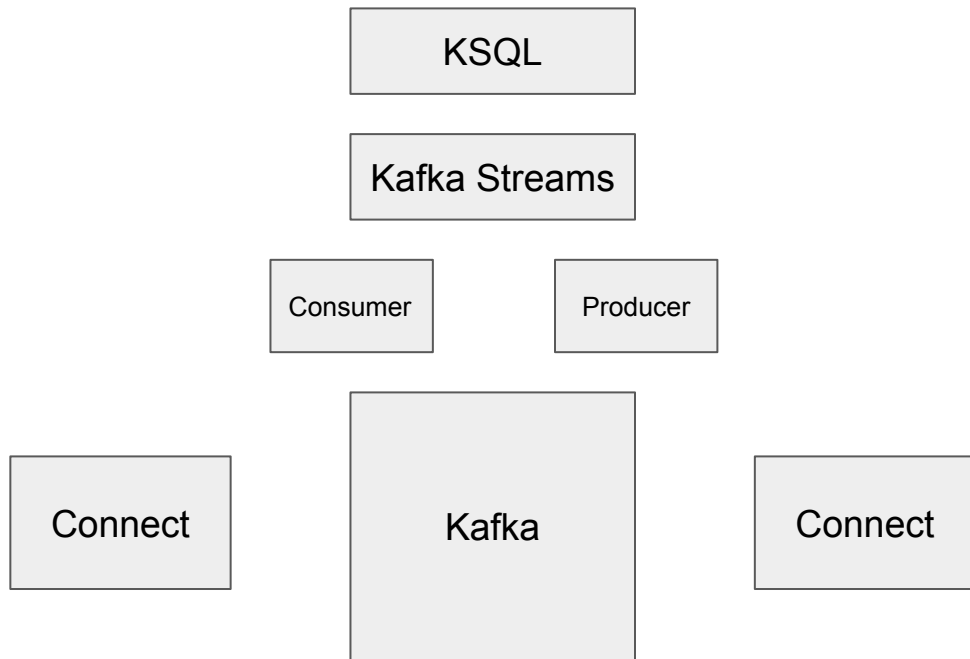
2019-01-16

Stream Processing

- Unbounded, continuous stream of events
 - Real-time vs. (micro-)batching
 - Handle high volume & velocity data
-
- Filtering & Enrichment
 - Validation & Transformation
 - Dehydration & Hydration
 - Routing, Splitting and Joining



Kafka Ecosystem





KSQL

- An engine for running streaming jobs on Kafka data
 - Queries are defined in **SQL**-like syntax and run within **KSQL-server** instance
 - Record at a time, real-time processing
 - Input from and output to Kafka topics
-
- First release in August 2018, developing rapidly.




KSQL

- Queries and posted to KSQL server over REST interface
- KSQL commandline client



Why KSQL?

- Convenient, no need to write application code
- No dependency on a specific language runtime
- Continuous queries operate without needing to run / schedule application runtime
- Very fast development cycles



Why Streams / Consumer/Producer?

- KSQL might not offer the level of controls than Streams / Consumer/Producer
- KSQL has access to Kafka topic data only, e.g. other systems
- KSQL is somewhat tricky to extend with custom functions
- KSQL created topologies tend to use a lot of intermediate topics; data reshuffling may amplify network utilization quite a bit



KSQL basics

Define a **stream** for input:

```
CREATE STREAM pageviews  
  (status INTEGER,  
   url STRING,  
   user STRING)  
WITH  
  (KAFKA_TOPIC='pageviews',  
   VALUE_FORMAT='JSON');
```




KSQL basics

Query a **stream**:

```
SELECT status, user, 'staging_' || url FROM pageviews  
WHERE status = 200;
```

Filter + transform

KSQL basics

Create a new **stream** & **topic**:

```
CREATE STREAM pageview_errors
  WITH
    (KAFKA_TOPIC = 'pageview_errors',
     VALUE_FORMAT='JSON')
  AS
    SELECT url, status, user
       FROM pageviews
      WHERE status <> 200
      PARTITION BY url;
```

Background query



KSQL basics

Check what's running:


```
SHOW STREAMS;
```

```
SHOW QUERIES;
```

```
DESCRIBE pageviews;
```

Streams & tables

Key	Value
a	blue
b	green
a	violet
b	yellow



Key	Value
a	violet
b	yellow





KSQL

Repartition:

```
CREATE STREAM pageviews_by_user  
  AS  
    SELECT user, url, status  
    FROM pageviews  
    PARTITION BY user;
```

KSQL

Aggregations:

```
CREATE TABLE pageviews_per_user
WITH
  (KAFKA_TOPIC = 'pageviews_per_user',
  VALUE_FORMAT='JSON')
AS
  SELECT user, count(*) as count
  FROM pageviews
  GROUP BY user;
```

Aggregation functions: max, min, count, ...

KSQL

Windowed aggregations:

```
CREATE TABLE pageviews_per_minute
WITH
  (KAFKA_TOPIC = 'pageviews_per_minute',
   VALUE_FORMAT='JSON')
AS
  SELECT url, count(*)
  FROM pageviews
  WINDOW TUMBLING (SIZE 1 MINUTE)
  GROUP BY url;
```

Windowing: tumbling, hopping, session

KSQL

Joins:

```
CREATE STREAM pageviews_with_user_stats
AS
  SELECT p.url, ppu.count
  FROM pageviews_by_user p
  LEFT JOIN pageviews_per_user ppu
    ON p.user = ppu.user;
```

Joins: left, inner, outer

stream - stream (windowed), stream - table (lookup)

stream-stream: event correlation/attribution, stream-table: data enrichment

KSQL

Joins:

```
CREATE STREAM feedback_with_pages
AS
    SELECT pbu.user, pbu.url, pbu.status
    FROM feedback_by_user fb
    INNER JOIN pageviews_by_user pbu
    WITHIN 1 MINUTE
    ON fb.user = pbu.user;
```

Joins: left, inner, outer

stream - stream (windowed), stream - table (lookup)

stream-stream: event correlation/attribution, stream-table: data enrichment

KSQL Benefits

- Quick to develop with - no code needed
- REST Can be called from any language
- Kafka Streams backs state & progress into Kafka
- Distributed operation: work distribution, fault recovery
- Scale by number of instances
- Trivially runs on cloud & containers



Links & resources

- <https://docs.confluent.io/current/ksql/docs/tutorials/>
- <https://github.com/confluentinc/ksql>
- <https://docs.confluent.io/current/ksql/docs/developer-guide/syntax-reference.html>



Epilog

- KSQL license changes / Dec 2019: no longer Open Source
- We're looking at Flink ourselves, and likely present our experiences in the coming meetups.



