



# Scaling Kafka in the Cloud

Meetup, Credorax Group

Heikki Nousiainen

2019-09-03

Tel Aviv

# Speaker

- Heikki Nousiainen
- CTO, co-founder @ Aiven
- First contact with Apache Kafka in 2014



@hnousiainen



# Aiven

- *Your data cloud*
- Based in Helsinki and Boston
- 8 data engines now available in 6 clouds and 80 regions, virtual and bare metal instances
- Launched a fully-managed Kafka Cloud service in 2016
- Other services: PostgreSQL, MySQL, Elasticsearch, Cassandra, Redis, InfluxDB



# Kafka @ Aiven

- Managed Kafka Offering since 2016
- 100s of Kafka Clusters
- 1000s of Brokers
- Median cluster 600 messages / second
- Median message size just over 200 bytes
- Typical cluster cost: \$660 / month



# Kafka @ Aiven

- We utilize Kafka as core message bus between all of our components
- Signaling, Metrics, Logs
  - Configuration/state communication
  - Logs from VM to Kafka, from Kafka to ES
  - Stats from VM to Kafka, from Kafka to InfluxDB / M3
- Why Kafka? Needed a fault tolerant signaling for Postgres failovers



# Our Challenges & Constraints

- Six supported cloud providers
  - Immutable infrastructure
  - Keeping software up-to-date
  - Adjustable cluster size and broker specs
  - Migration between regions and clouds
  - High availability, SLAs
- 
- Changing instances, IP addresses
  - Ephemeral disks
  - Roll-forward updates





# Aiven Architecture



# Aiven Architecture



- Kafka Brokers implemented as dedicated VMs
- Backed by management layer responsible for provisioning the cluster resources







# Aiven Architecture

- Management layer provisions and provides resources to the cluster.
- Management layer monitors and ensures cluster meets the specification: number of brokers, size of brokers, correct software version, broker health
- Availability Zones for fault tolerance



# Aiven Architecture

- Kafka Brokers managed by agent software running within VM
- Per cluster ZK co-hosted on same VMs
- Agent responsible for setting up and managing both ZK and Kafka
- Immutable; software updates via VM replacement



# Aiven Architecture

- Intra-cluster IPv6 overlay network
- Built using IPSec tunnels over either public or private instance IP addresses
- Consistent addressing regardless of the VM location / networking capabilities



# Building Blocks



# ZooKeeper

- Distributed, consensus based synchronization & coordination service
- Used by Kafka for controller election and cluster state management
- ZooKeeper 3.5 branch supports dynamic reconfiguration
- Released May 20th, 2019
- In use in Aiven for 3+ years



# ZooKeeper

- Dynamic reconfig requires quorum
- Fault recovery adds new node(s), removes disappeared ones
- Roll-forward upgrade adds new node(s), removes obsoleted nodes
- Both failure and positive cases handled using the same procedure




# Kafka Configuration Changes

- There are more and more settings that can be configured dynamically
- Some Kafka settings require broker restart
  - `auto.create.topics.enable`
  - `offsets.retention.minutes`
- Agent restarts Kafka process as needed
- Interlocking / synchronization via ZK
  - Avoid restarting multiple brokers at once



# Kafka Corrective Actions - Broker

- Kafka broker sometimes falls off sync with the cluster
- Also, thread deadlocks
- Restart forces full state reload
- Can be triggered automatically
- For availability, must not to restart more than one broker at a time





# Active Partition Placement

- Agent performing continuous monitoring for partition placement
- Adjusts placement automatically:
  - To satisfy redundancy requirements & data distribution
  - To balance storage usage
  - To balance partition leaders

Process:

- Create and maintain a continuously updated plan in ZK
- Correct deviations from the plan
- Throttle the amount of inflight partition movements
- Partition management instructions via ZK

`/admin/reassign_partitions`



# Partition Placement - Cluster Changes

- Drain nodes that are about to be decommissioned
  - Maintain the specified level of redundancy
  - Maintain correct AZ spread
- Ensure clients get negative reply  
(not\_leader\_for\_partition) and reload metadata
- Our placement logic is part of our per VM management agent
- My hope is to separate it & publish it as Open Source



# Kafka Corrective Actions - Cluster

- Management services detect loss of Kafka brokers
  - Creates new VM resource for the cluster
  - Assigns a new Broker ID
- 
- Agent adjusts the ZK cluster memberships
  - Partition placement uses the new resource and restores replication levels

# Kafka Cluster Operations

- Management services detect:
  - Required security software updates
  - Requested Kafka version updates
  - Mismatching VM types - scaling
  - Mismatching VM number - scaling
  - Mismatching VM locations - migration
- Create a single new VM
- Partition manager drains an existing broker - data migration
- Old broker recycled
- Loop until complete



# Chaos Engineering

- Practice makes perfect
- Kafka is critical component in our software stack
- We initially run our Kafka clusters using Google Cloud Platform preemptible VM instance types
- The instances can be terminated at any time, and will be terminated after 24 hours




# Monitoring




# Monitoring

- We monitor Kafka clusters by reading JMX via Jolokia
- Collected by telegraf and sent to centralized store
- Key metrics for alerts:
  - UnderReplicatedPartitions (\*)
  - IsrShrinksPerSec
  - RequestHandlerAvgIdlePercent / NetworkProcessorAvgIdlePercent
  - LogCleanerManager:time-since-last-run-ms
- If you have stable workloads, consider low/high message rates
- Implement paging and/or automatic corrective responses



# Monitoring

- Metrics generated from Kafka log / journald
  - Exceptions and Errors
    - ReplicaFetcher errors
    - LeaderAndIsr requests
  - Can be noisy!
- 
- We're looking to contribute in this area






# Monitoring

- ZooKeeper
  - `/brokers/topics/...`
  - `/admin/reassign_partitions`
  - `/admin/delete_topics/` - if you're using ZK for topic management
- Health checks to API
  - metadata call
  - response times



# Partition Placement - Monitoring

- Monitor progress via `/admin/reassign_partitions`
- Sometimes gets stuck:
  - Ever tries to implement reassignment on a broker that's gone
  - Failed partition leadership assignments
- On timeout, force controller re-election
- Hoping to investigate and contribute fixes in Kafka
- Beware of UnderReplicatedPartitions
- We filter UnderReplicatedPartitions vs.  $ISR \geq \text{number of desired replicas}$



# Thank you!

@hnousiainen  
htn@aiven.io

<https://aiven.io>

