



Scaling Kafka in the Cloud

All Things Open

Heikki Nousiainen

@hnousiainen

2019-10-14

Raleigh, NC

Speaker

- Heikki Nousiainen
- CTO, co-founder @ Aiven
- First contact with Apache Kafka in 2014



@hnousiainen



Aiven

- Based in Helsinki, Boston and Sydney
- 7 data engines now available in 6 clouds and 80 regions, virtual and bare metal instances
- Fully-managed Kafka Cloud service launched in 2016
- PostgreSQL, MySQL, Elasticsearch, Cassandra, Redis, InfluxDB



Kafka @ Aiven

- Managed Kafka Offering since 2016
- 500 Kafka Clusters under management
- Cluster sizes range from 3 to 30 brokers
- Over 10,000 services in total under management



Kafka @ Aiven

- We utilize Kafka as core message bus between all of our components
- Signaling, Metrics, Logs, Audit logs
 - Configuration/state communication
 - Logs from VMs to Kafka, from Kafka to Elasticsearch
 - Stats from VMs to Kafka, from Kafka to InfluxDB / M3



Kafka @ Aiven

- Decoupled architecture
 - Handle various kinds of network splits
 - Avoid blocking / waits on the producer side
- Scale
 - Kafka scales near linear by number of brokers
 - Adjustable number of consumer processes
- One interface, multiple use cases
- One set of data, multiple use cases
- Security: ACL controls over read/write access





Aiven Architecture - principles

- Everything is code. Aiven is a software product.
- Automation to handle full lifecycle operations: deployment, operations and fault response.
- Converging operation: code works towards to the desired service state until established.
- Immutable infrastructure: ensure known state by enforcing rolling upgrades.





Aiven Architecture

- Kafka brokers implemented as dedicated VMs
- Backed by management layer responsible for provisioning the cluster resources





Aiven Architecture

- Cluster members managed by agent software running within VM
- Per cluster ZK co-hosted on VMs
- Agent responsible for setting up and managing both ZK and Kafka
- Immutable; software updates via VM replacement



Aiven Architecture

- Management layer provisions and provides resources to the cluster.
- Management layer monitors and ensures cluster meets the specification: number of brokers, size of brokers, correct software version, broker health
- Availability Zones for fault tolerance



Our Challenges & Constraints

- Six supported cloud providers
 - Immutable infrastructure, roll-forward updates
 - Keeping software up-to-date
 - Adjustable cluster size and broker specs
 - Migration between regions and clouds
 - High availability, SLAs
-
- Changing instances, IP addresses
 - Ephemeral disks



Networking

- Span of private networks vary between clouds.
 - Migration across networks can be tricky: mixed private & public addressing.
 - Kafka advertises addresses in protocol endpoints also for interbroker communication.
-
- We utilize IPv6 overlay network between cluster nodes
 - Consistent addressing regardless of the location or network



Networking

- Point-to-point IPv6 tunnels over IPv4 IPsec connections
 - Libreswan
 - Managed with whack commandline client
 - Dynamically select the best routing option between any two VMs
-
- All internode communication authenticated and encrypted
-
- Closely tracking Wireguard progress as well



ZooKeeper

- Distributed, consensus based synchronization & coordination service
- Used by Kafka for controller election and cluster state management
- Changes in ZK cluster used to require restarts
- ZooKeeper 3.5 supports dynamic reconfiguration
- Add/remove nodes, switch roles
- Released May 20th, 2019
- In use at Aiven for 3+ years




ZooKeeper

- Dynamic reconfig requires quorum
- Fault recovery adds new node(s), removes disappeared ones
- Roll-forward upgrade adds new node(s), removes obsoleted nodes
- Both failure and positive cases handled using the same procedure



Controlled Kafka Broker restarts

- There are more and more settings that can be configured dynamically
- Some Kafka settings require broker restart
 - `inter.broker.protocol.version`
 - `auto.create.topics.enable`
 - `offsets.retention.minutes`
- For availability, avoid restarting all brokers at the same time



Controlled Kafka Broker restarts

- Agent detects a change requiring broker restart
- ZooKeeper based locks ensure that a single broker restart is initiated and completed at any one time



Controlled Kafka Broker restarts

- Same protocol applies to automated error recovery
- Handle Kafka Broker misbehaviour:
 - Deadlocks
 - Memory leaks / thread exhaustion
 - ZK induced cluster state mismatches
- Detect and correct automatically

Active Partition Management

- Agents perform continuous monitoring of partition placement
- Automatic adjustment for:
 - Satisfying number of alive replicas
 - Satisfying data replication across zones
 - Balance storage usage
 - Balance partition leader assignments
- Controlled draining of decommissioned brokers
 - Maintain replication level, AZ spread
 - Ensure clients receive “Not Leader For Partition”



Active Partition Management

- Calculate and maintain optimal plan in ZK
- Detect deviations from the plan
- Select and request a set of partition reassignments
 - `/admin/reassign_partitions`
 - Throttle amount of ongoing requests
 - Optimize & balance broker to broker links
- Currently coupled with our agent logic
- I'm looking to split & Open Source as a standalone daemon




Kafka Cluster operations

- Loss of Kafka Broker is detected by management plane
- New VM resource is created and joined into the cluster created
- A new Broker ID assigned
- Agent adjusts the ZK cluster memberships
- Partition placement uses the new resource and restores replication levels



Kafka Cluster operations

- Management plane detects:
 - Required security software updates
 - Requested Kafka version updates
 - Mismatching VM types - scaling
 - Mismatching VM number - scaling
 - Mismatching VM locations - migration
- Create new VMs
- Partition manager drains outgoing brokers - data migration
- Old brokers recycled



Monitoring

- JMX metrics via Jolokia
 - ZooKeeper keys and values
 - Logs
 - Health checks, metadata calls
-
- Utilized by the local agent for decision making and corrective actions
 - Streamed with Telegraf to central processing using Kafka



Monitoring

- JMX metrics via Jolokia
 - UnderReplicatedPartitions (*)
 - IsrShrinksPerSec
 - RequestHandlerAvgIdlePercent /
NetworkProcessorAvgIdlePercent
 - LogCleanerManager:time-since-last-run-ms
- For topics / partitions with stable workloads, monitor & alert on high/low messages / sec bands



Monitoring

- ZooKeeper keys and values
 - /admin/reassign_partitions
 - /brokers/topics/...
- E.g. /admin/reassign_partitions update age
- Sometimes Kafka gets stuck:
 - Ever tries to implement reassignment on a broker that's gone forever
 - Failed partition leadership assignments
- On timeout, trigger Kafka controller re-election



Monitoring

- Logs
 - Exceptions & Errors
 - ReplicaFetchError
 - LeaderAndIsr requests
 - Log Truncations
 - Derived metrics, errors / minute
- Can be quite noisy
- Local decisions, alerts on metrics, centralized collection
- Journald to Kafka to Elasticsearch



Keeping up with Kafka versions

- We've used Kafka since version 0.8.2
- Extremely good backwards compatibility support
- Quality steadily improving:
 - Memory Leaks
 - Resource starvation
 - Deadlocks
 - Thread exhaustion (Exceptions)
 - Falling out of sync with the cluster
 - 2.1 => TLA+ modeled replication protocol improvements
- Generally a lot less broker restarts / forced controller re-elections

Break Things!

- Kafka is critical component to our infrastructure
- Practice makes perfect
- Terminate broker & ZK processes & VMs
- Induce network splits and errors
- Destroy volumes
- Spot & pre-emptible instance types
- Verify availability and capability to operate clusters



Thank you!

@hnousiainen
htn@aiven.io

<https://aiven.io>

