



Apache Kafka 101: The Why, What, and How



John Hammink
Developer Advocate



Chris Gwilliams
Solutions Architect

Aiven brings the best Open Source data technologies to all public clouds



Kafka



PostgreSQL



Elasticsearch



Cassandra



Grafana



Redis



InfluxDB



MySQL

Coming soon...



Clickhouse



M3

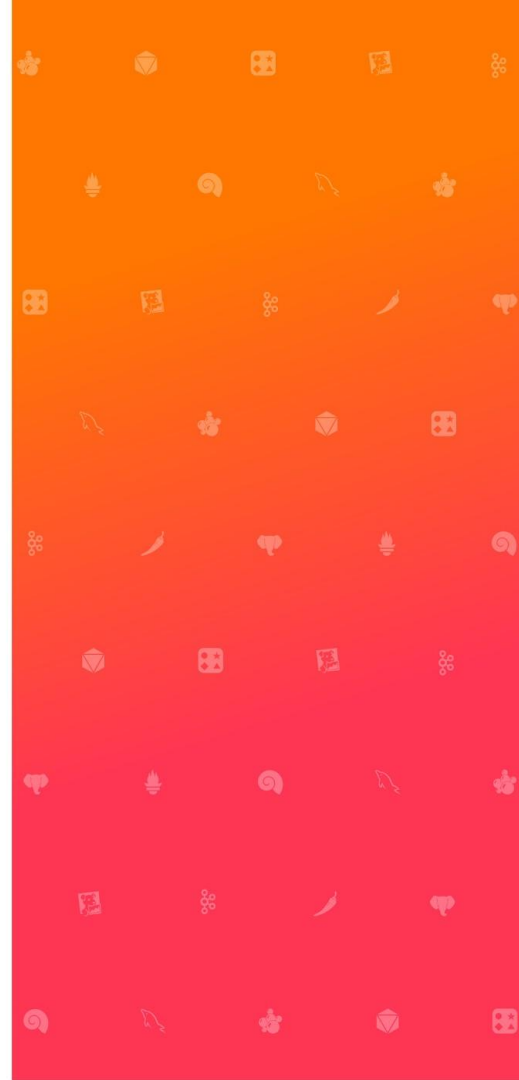


Flink



Agenda

- Introduction
- Kafka -
 - WHY?
 - WHAT?
 - HOW?
- Kafka ecosystem
- DEMO
- Q & A (aka HUH?)

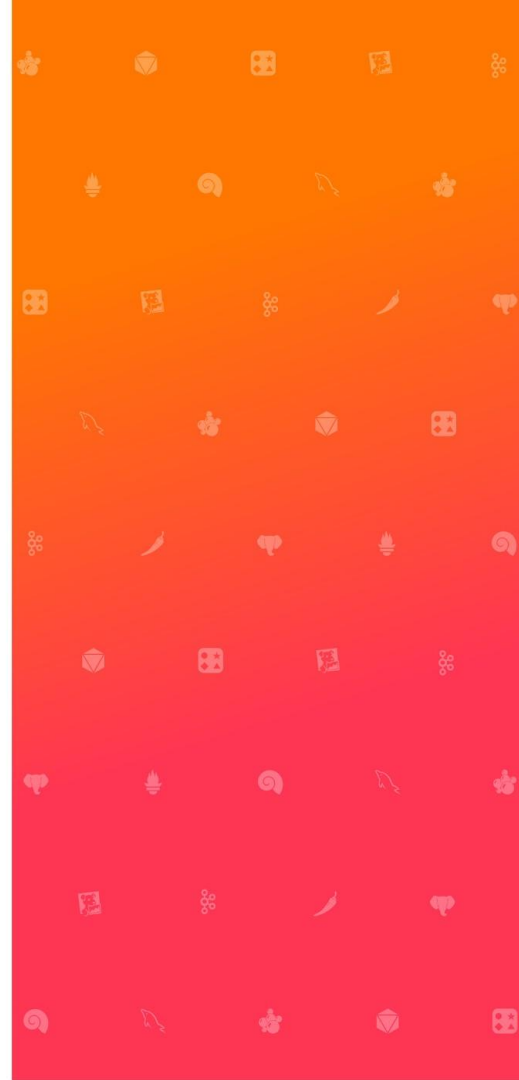


Background: the WHY of Kafka



History of Apache Kafka

- shift away from “monolithic” programs to those that must pass data between ***senders*** and ***receivers***
- eventually, systems became **distributed**
- imagine the complexity with **millions of nodes** sending items back and forth like **status updates**, **presence**, **login states**, among other things. Imagine the volume of messages here.
- complexity = $n * n$
- **Also, decoupling...**

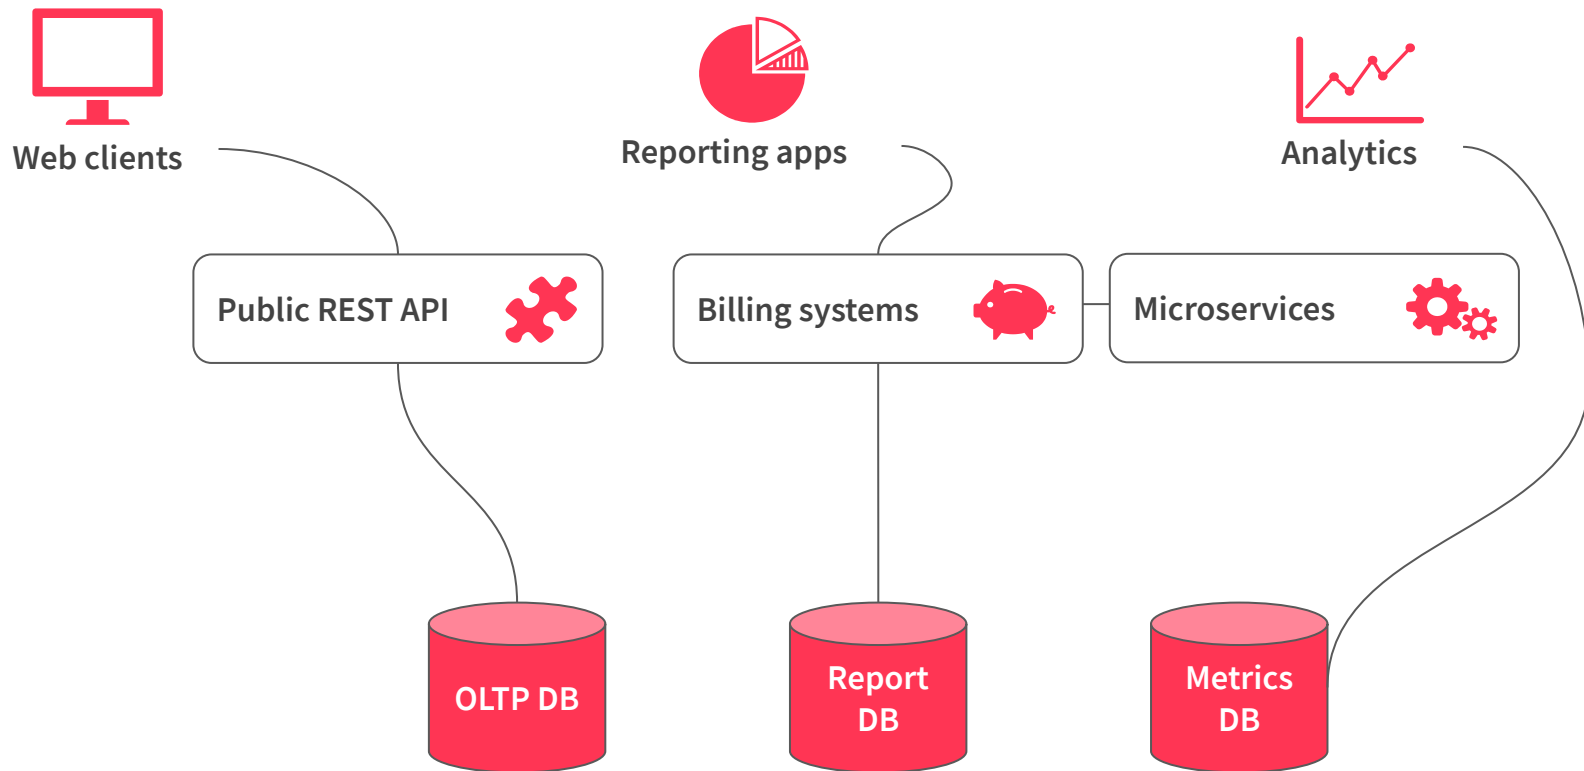


History, part II

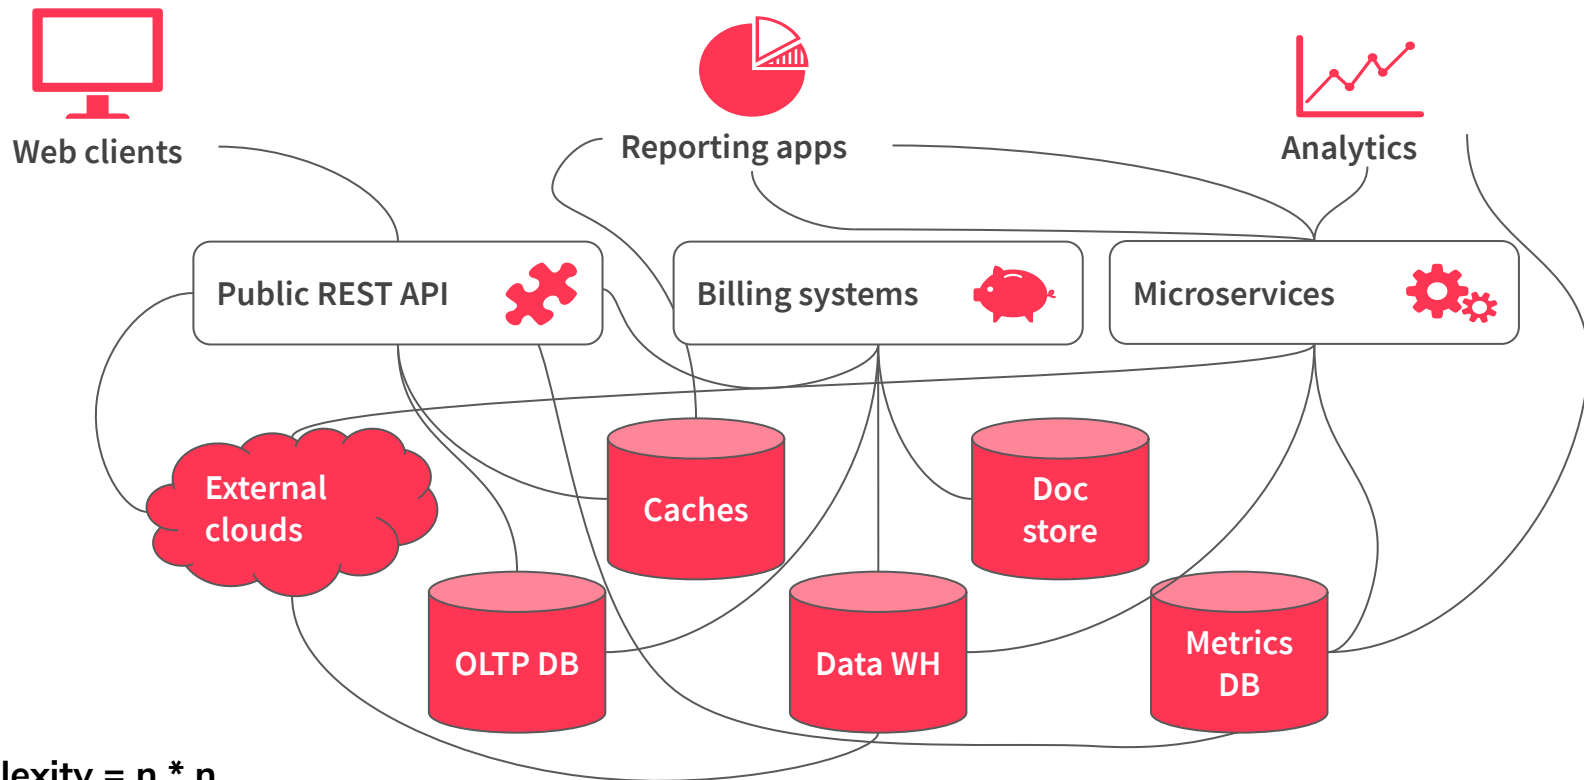
Alas, a middle step:



“Traditional” data flow model

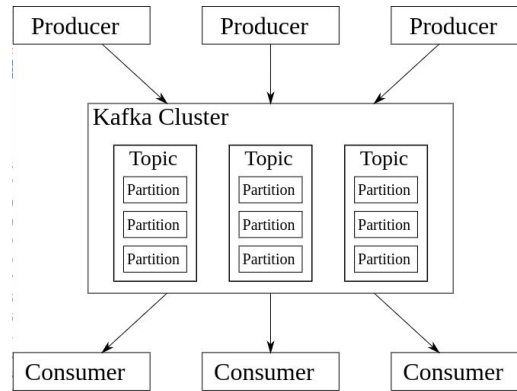


“Traditional” data flow model

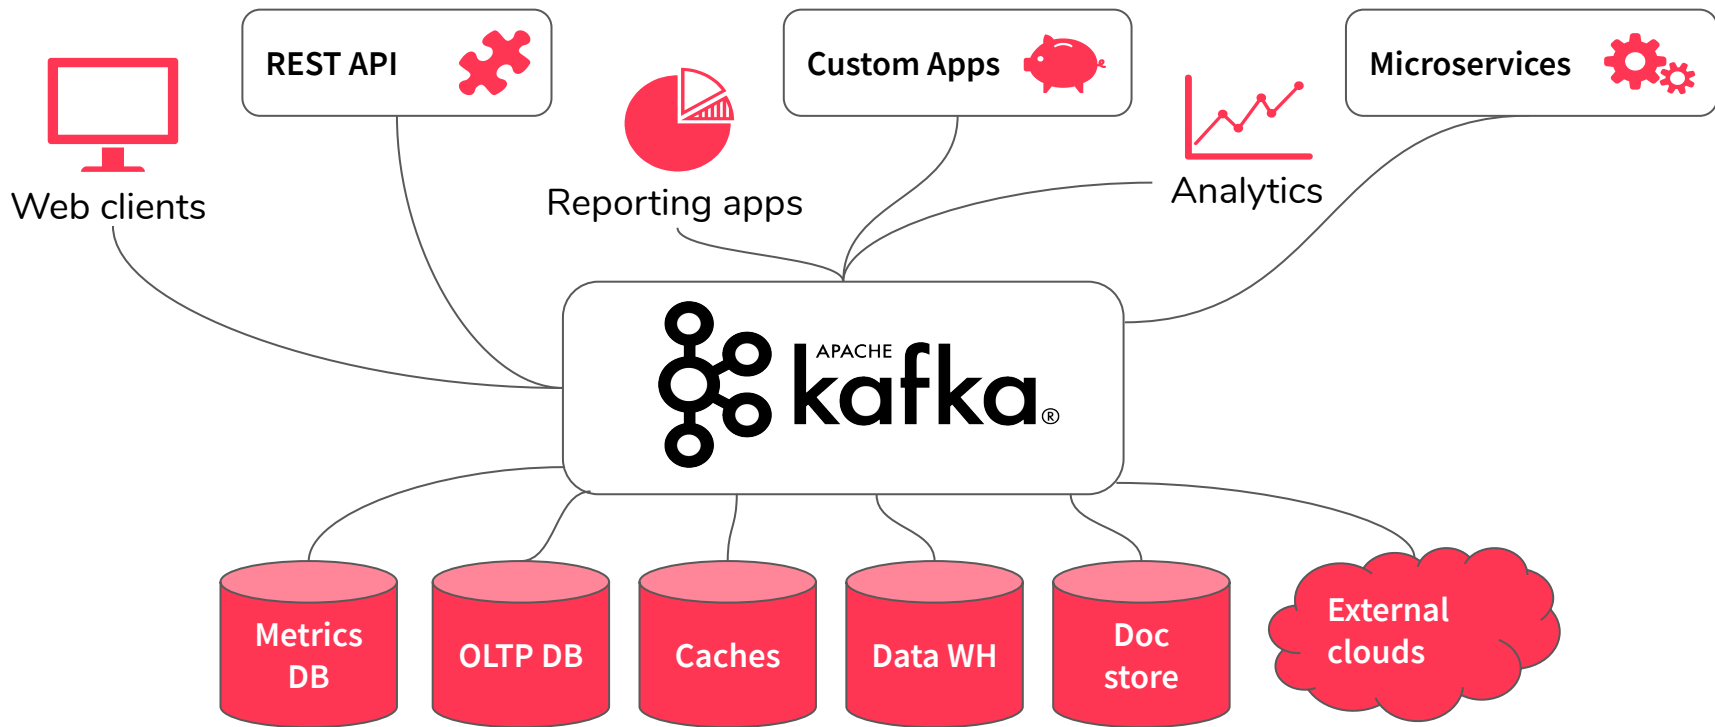


Enter: Apache Kafka!

- Developed at LinkedIn starting in 2010
- open-sourced to Apache in early 2011
- improved on the pub/sub architecture
- complexity = $n + n$



Kafka-centric data flow model

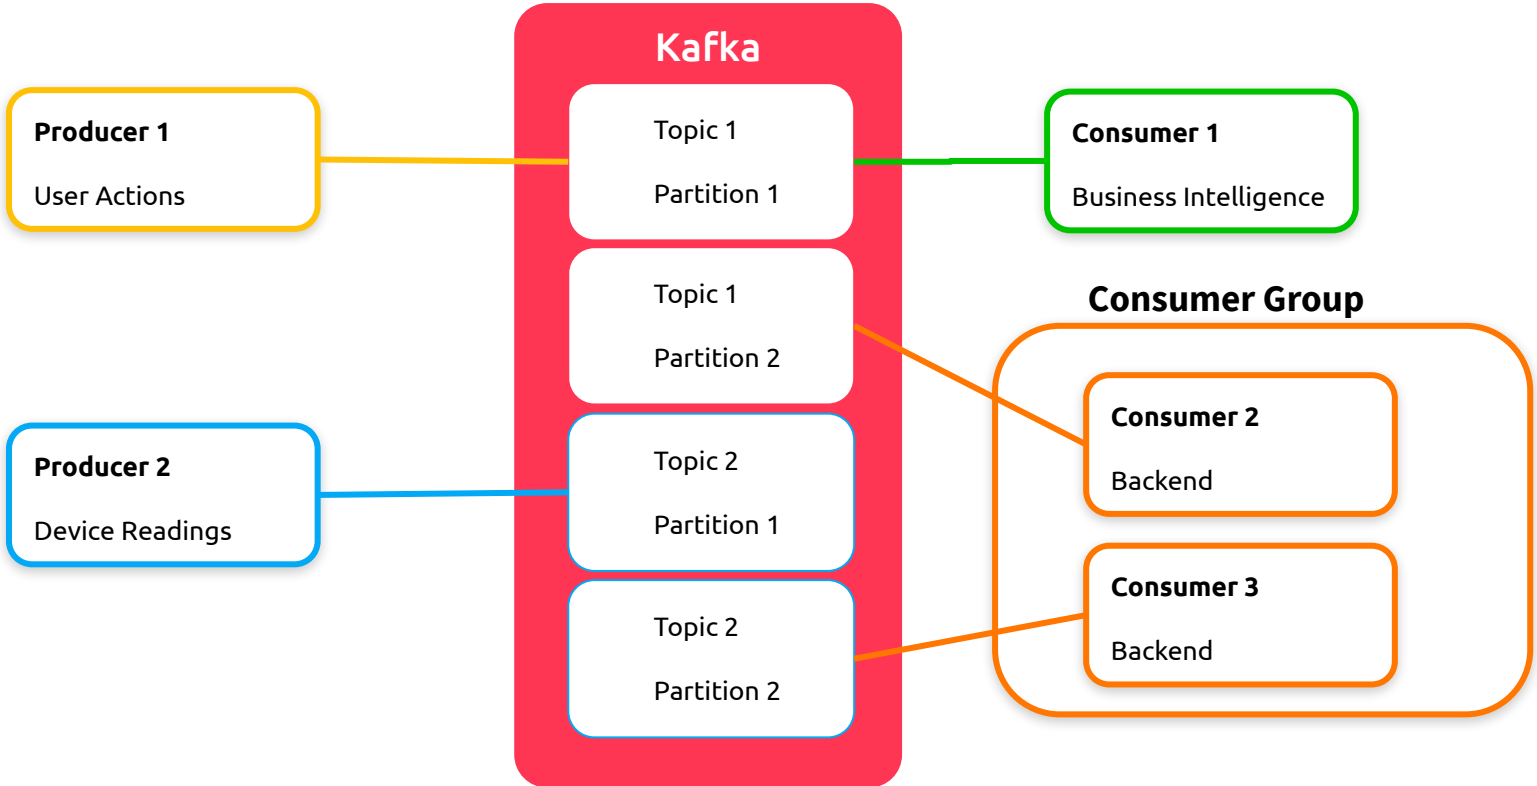


complexity = n + n

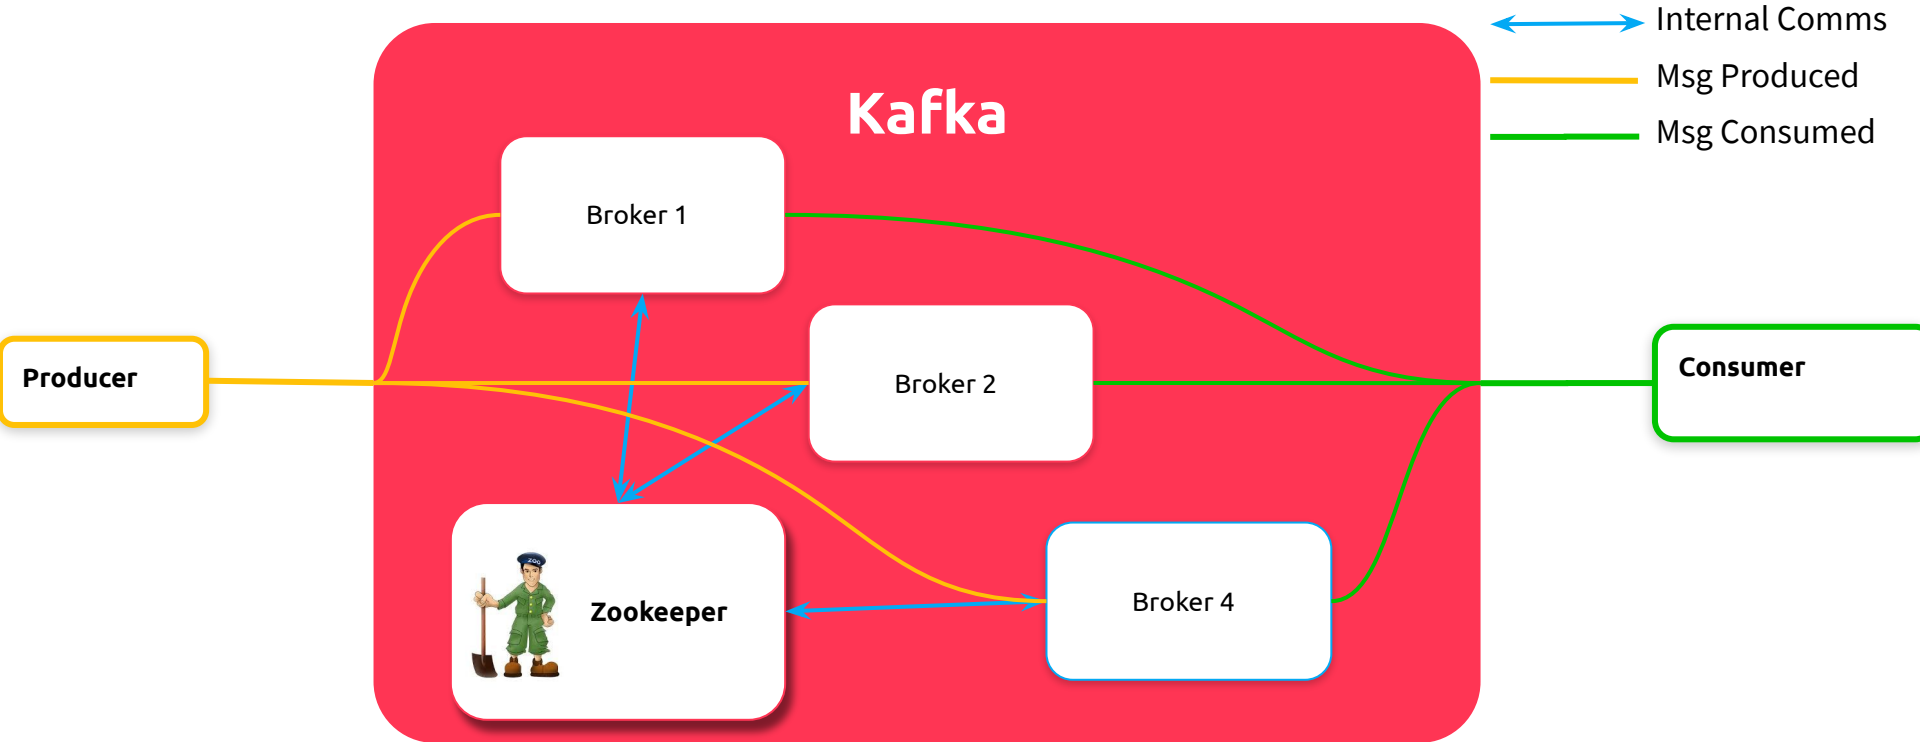
Architecture: the WHAT of Kafka



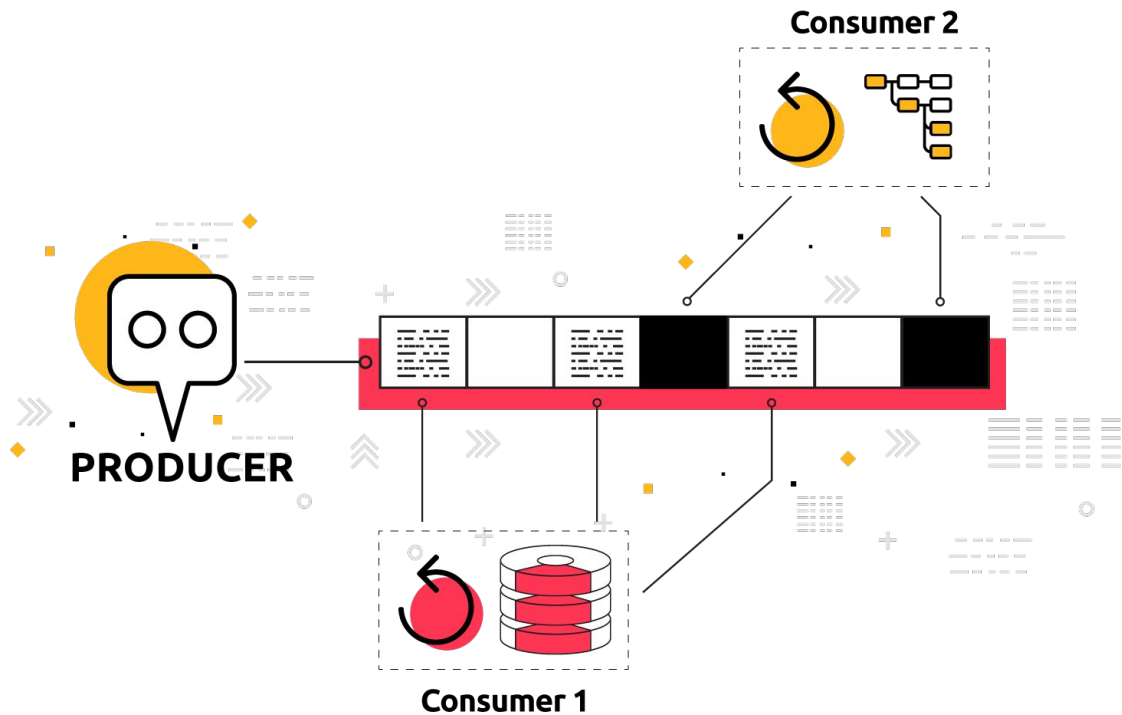
Producin' and Consumin'



Brokers, Clusters, ISR and Zookeeper



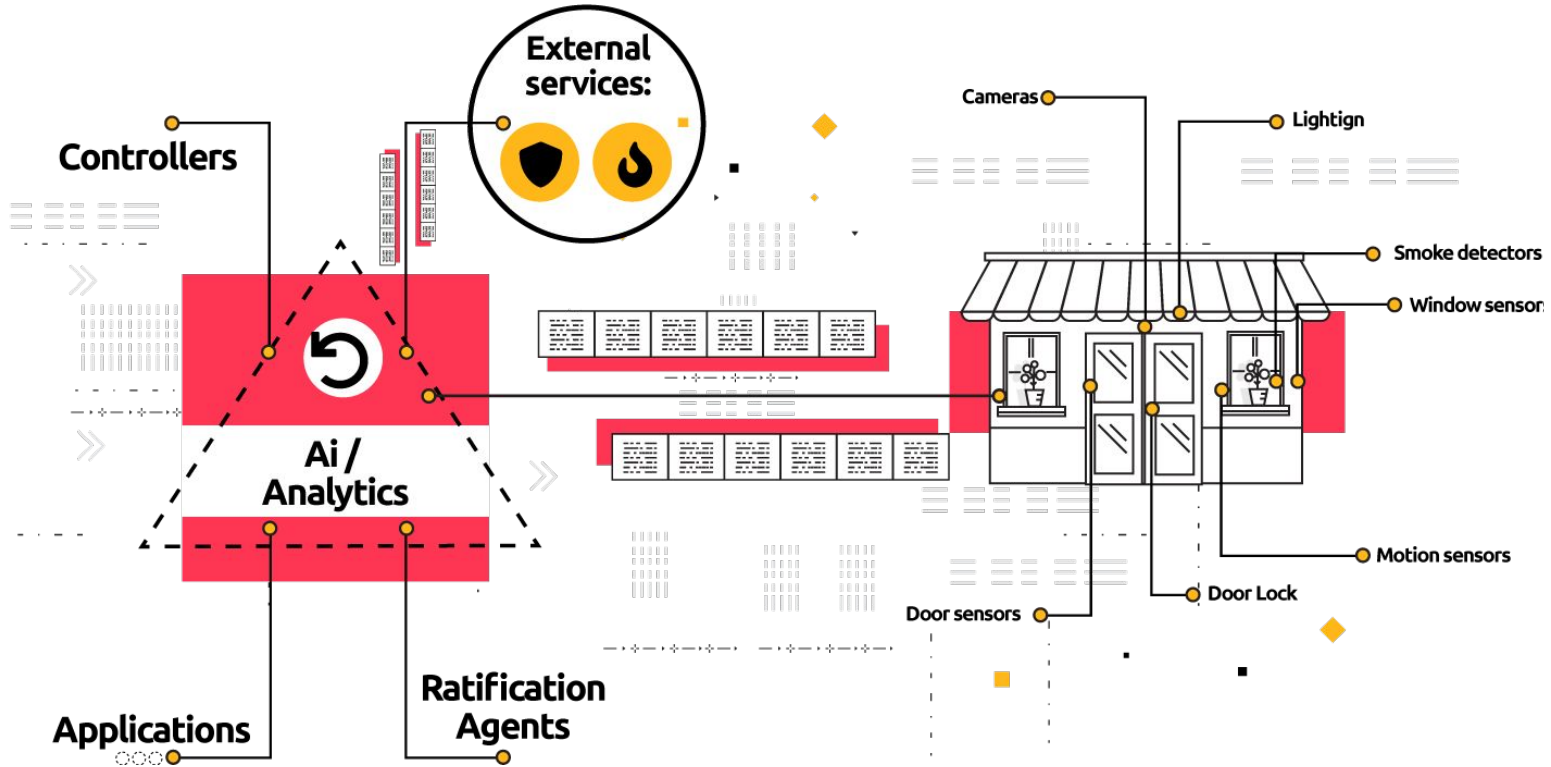
What Apache Kafka does: Publish/Subscribe



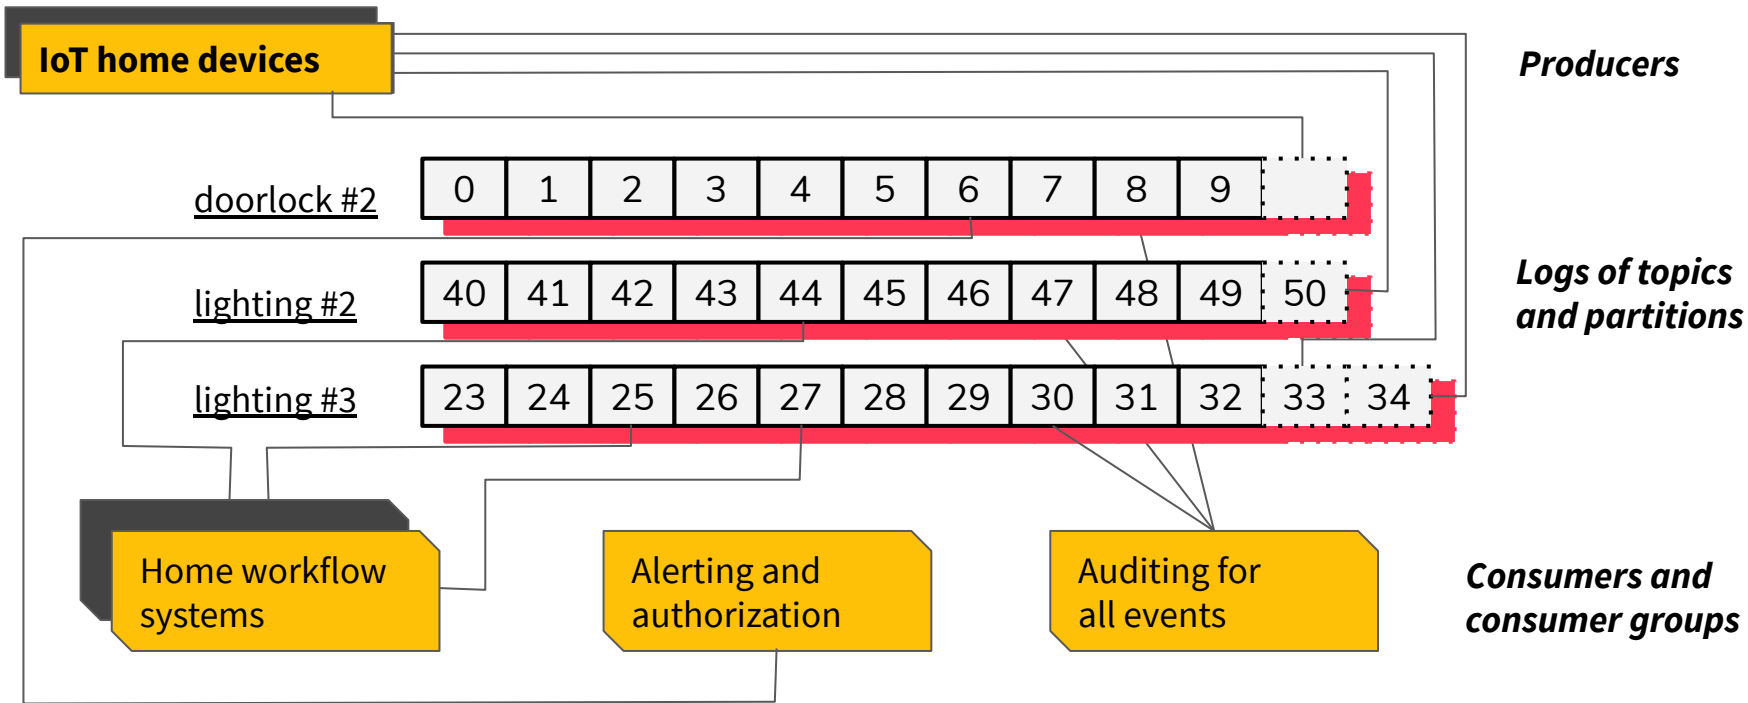
Use-Cases: the HOW of Kafka



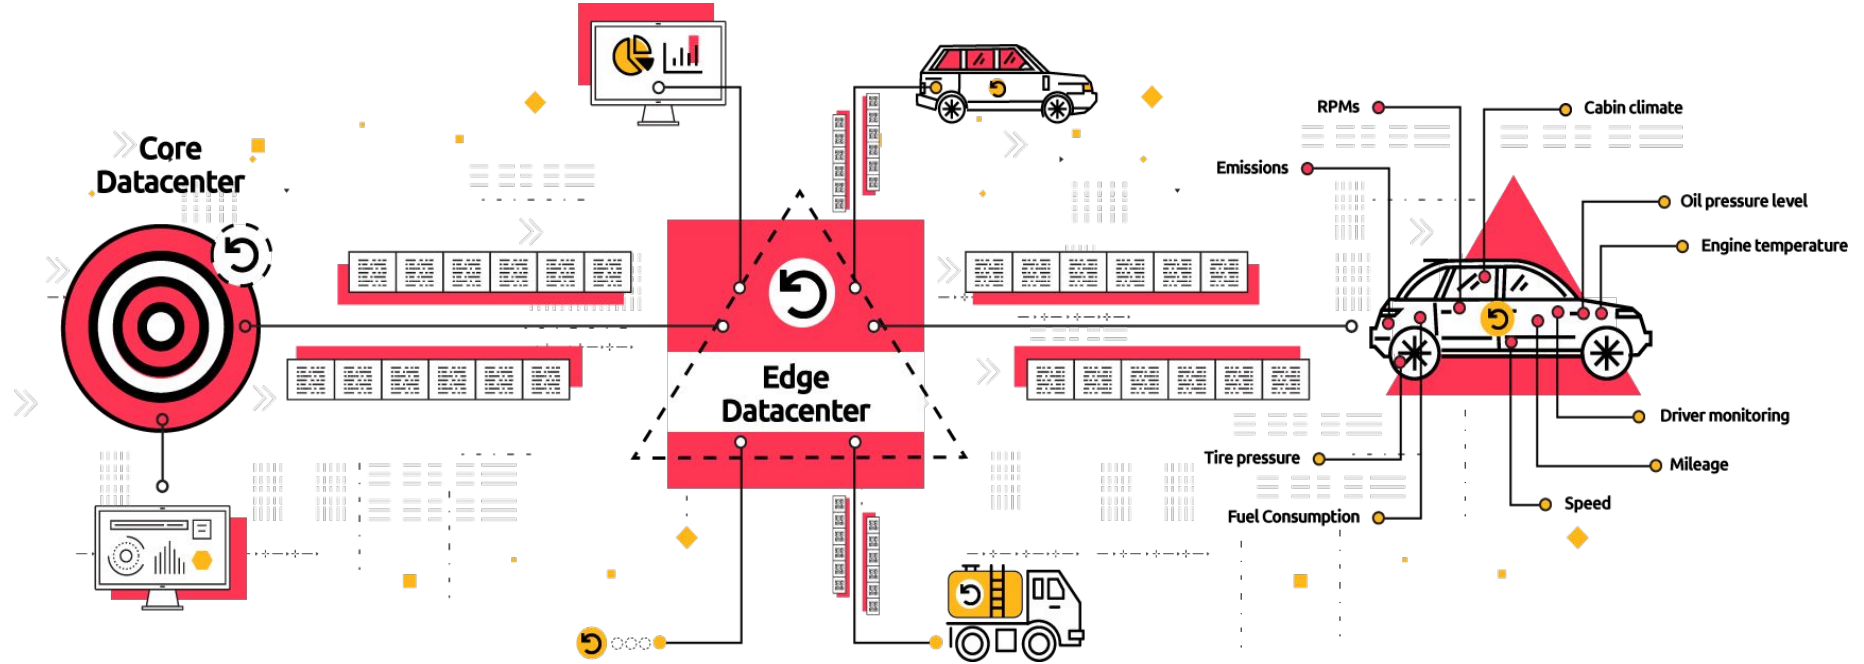
IoT event processing example



How Kafka handles that



Core to Edge example



Advantages of Managed Kafka



WHY Apache Kafka?

- Routing rules configured by consumer
- Built-in support for asynchronous messaging
- High throughput/low latency



Operating Apache Kafka

Covers data pipeline requirements

- Scales to billions of messages per day
- Supports rack and data center aware replication
- Cross-region replication using MirrorMaker
- Real-time streaming
- Decoupling of message consumption & producing
- Client libraries & tools available for all popular languages

But maintenance can be a burden

- Depends on ZooKeeper
- Rebalancing of leaders and partitions
 - In case of failure
 - When scaling up or down
- Broker hangs
- Consider using a managed Kafka service, available from multiple vendors including us



WHY Managed Apache Kafka?

- Easy deployment
- No outlay costs
- Thousands of hours of experience
- Auto-scaling
- Centralized, no-fuss management
- High Availability

Managed is Peace of Mind

Focus on your product and let us take care of your infrastructure

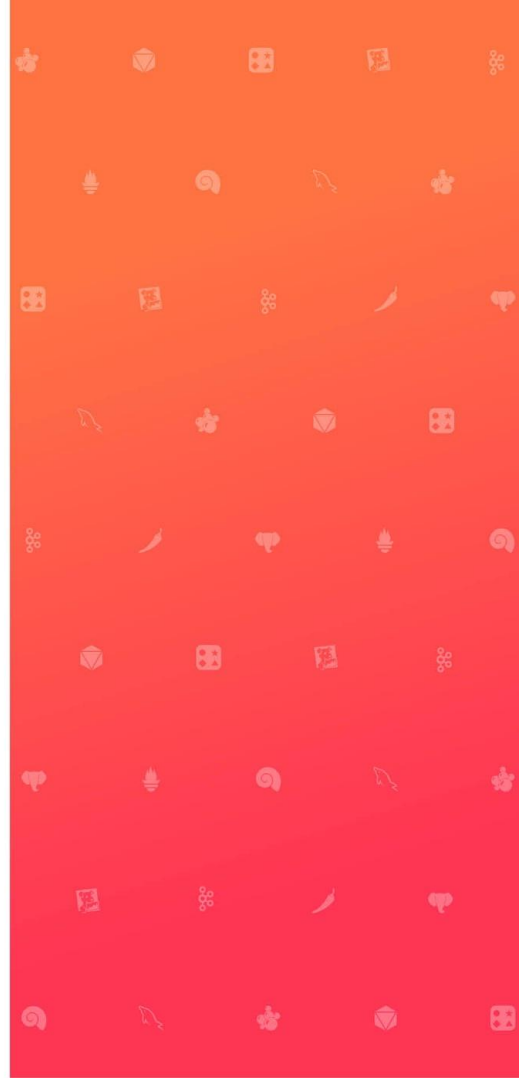


The Kafka Ecosystem



The Kafka Ecosystem

- Kafka Connect
- Schema Registry
- Data Streaming
 - Kafka Streams
 - Spark/Flink



Kafka Connect

- Native Kafka framework for integrating multiple data sources
 - Source Connectors push data to Kafka
 - Sinks dump data
- Runs alongside Kafka OR separately
- Configure with JSON
- Available for: RabbitMQ, Google BigQuery, PostgreSQL and more



Kafka



PostgreSQL



Elasticsearch



Cassandra



BigQuery



Redis



InfluxDB



MySQL

 RabbitMQ



amazon
S3

Schema Registry

Validate

Ensure message
consistency

Validate message source

Many Formats

JSON

Avro

PROTOBUF

Compatibility

Per Schema compatibility

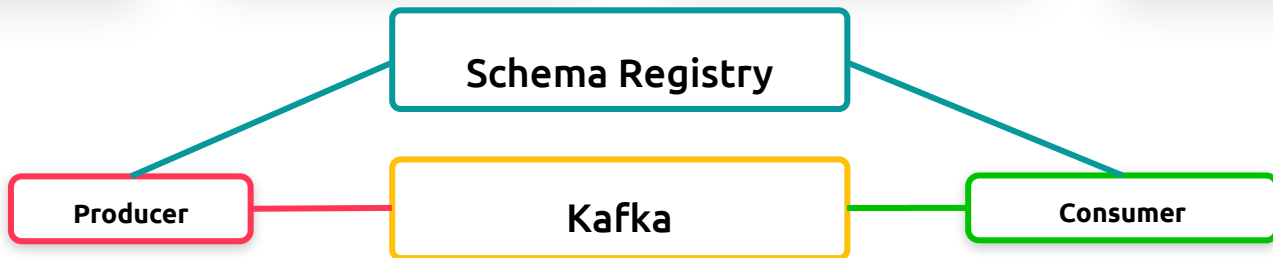
Backwards

Forwards

High support & service level

24/7/365 monitoring by
experienced Aiven
personnel

Dedicated phone support
available

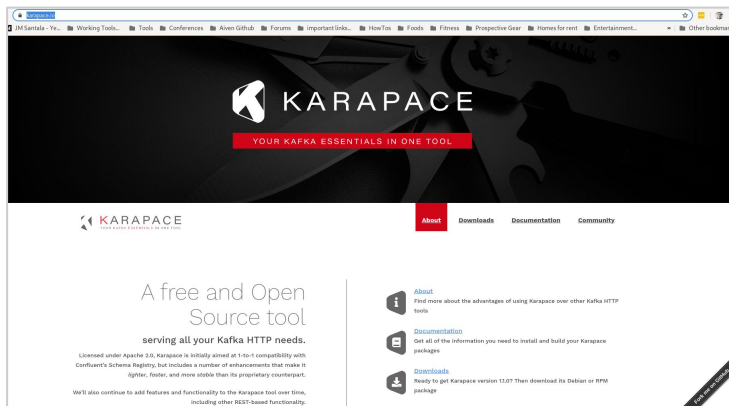


Karapace

github.com/aiven/karapace



aiven



Schema Registry++

OSI approved license

Use with ANY Kafka

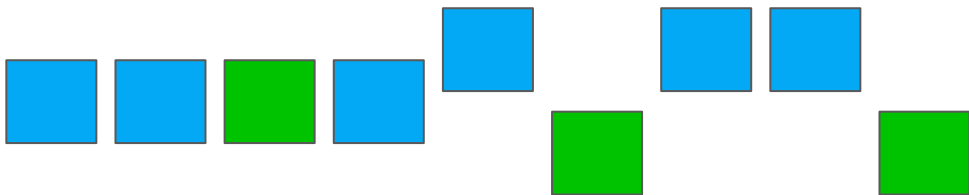
Lighter, faster, better, YOURS

Data Streaming -> Flink

- **Stateful Computations over Data Streams**
 - Flink lets you run stream and batch processing over any supported data source
- Processing can take place as SQL or as custom code running against the data
- Custom code written in Java or Python



Flink



Kafka Streams

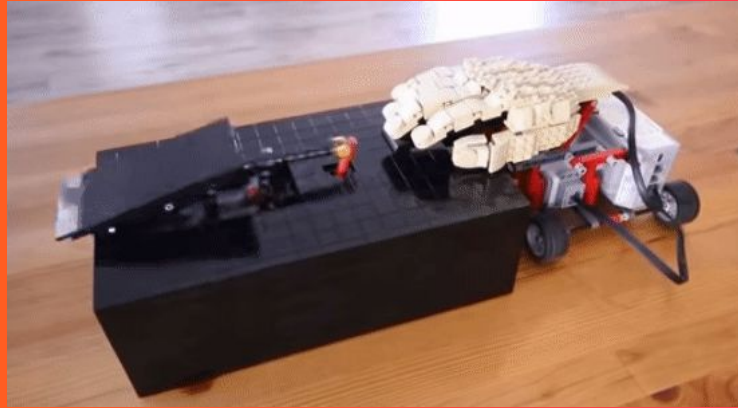
- Stream processing without the need for Spark/Flink
 - Java/Scala native (libraries for many other languages)
- Stateful
- Support out of the box
 - Essentially a Consumer

```
app = faust.App('myapp', broker='kafka://localhost')

# Models describe how messages are serialized:
# {"account_id": "3fae-...", amount": 3}
class Order(faust.Record):
    account_id: str
    amount: int

@app.agent(value_type=Order)
async def order(orders):
    async for order in orders:
        # process infinite stream of orders.
        print(f'Order for {order.account_id}: {order.amount}')
```

DEMO



Try it yourself!

https://github.com/aiven/presentations/tree/master/Cloud_chats_by_Aiven/Kafka_101

Final words

- **Kafka is:**
 - The evolution of messaging systems
 - The backbone of a Data Pipeline
 - An ecosystem of tools
 - Easy to integrate with popular software
 - Waiting for **you** at aiven.io

Produce once, consume everywhere



**Thank you for your
time!**



Q & A

Do You Have Any Questions?

[@aiven_io](#)

webinars@aiven.io



Links / Further Reading

- [Aiven Kafka](#)
- [Karapace](#)
- [Aiven's Presentations](#)
- [Kafdrop](#)
- [Using Kafdrop Web- UI with Aiven Kafka](#)
- [Effective Kafka](#) (ebook) (unaffiliated)